

Support-gated selective recovery under controlled multi-sensor corruption

Riyang Luo^{1,*}

¹School of Mechanical and Power Engineering, Nanjing Tech University,
30 South Puzhu Road, Nanjing 211816, Jiangsu Province, China

*Correspondence: lry@njtech.edu.cn

Abstract

Recovering from sensor corruption requires choosing between a stable base classifier and a recovery endpoint that can both correct and damage decisions. We formulate this choice as controlled-corruption selective recovery. Endpoints and routing thresholds were fitted before evaluation on fixed future batches. The strict test set contained 1,436 observations per mechanism for which a controlled fault reached an available sensor group. The frozen conditional router prevented 94.7% of recovery-induced negative transfers while retaining 9.7% of available corrections. Only five of ten fitted endpoint pairs passed a deployment-support gate proposed after the frozen experiment, evaluated retrospectively and intended for prospective use. Applying that gate retrospectively reverted unsupported pairs to the base endpoint and yielded 96.2% prevention, 6.6% retention and a mean area under the receiver-operating-characteristic curve of 0.8132, compared with 0.8115 for the base endpoint. Under equal correction and harm values, calibration-selected all-row and two-stage analytical controls yielded mean test utilities of 107.9 and 110.6 units per 10,000 observations; unconditional recovery yielded 132.1. On this fixed test set and with no compute penalty, the two-stage policy had the highest mean utility among the tested policies at a two-to-one harm ratio, but its 95% fitted-pair interval crossed zero. At ratios of five or ten every learned policy had negative mean test utility and the base endpoint dominated. Severity, sensor-group, mechanism and batch transport changed retention and utility materially. The evidence supports a reproducible risk-control protocol for controlled corruption, not natural-failure recovery or device-independent field performance.

Keywords: multi-sensor fusion; sensor corruption; fault-tolerant classification; selective recovery; negative transfer; cross-fitting

1 Introduction

Scientific and industrial systems increasingly combine several channels whose availability and fidelity can change after deployment. Missing inputs, delayed diagnostics and distribution shifts can invalidate a model even when its architecture performs well on a static benchmark [1–4]. This problem is distinct from ordinary random missingness because a faulty sensor can remain present and produce plausible but biased, clipped, noisy or unresponsive measurements. A robust fusion rule must therefore decide not only whether a channel exists, but how strongly it should influence the prediction.

Several training strategies improve tolerance to absent or inconsistent sources. Modality dropout, masking and self-distillation expose models to missing inputs [5–7]. Representation refinement, gradient decoupling and expert routing address source imbalance or train–test mismatch [8–11]. Reliability-aware fusion additionally uses predictive uncertainty, learned credibility

or rejection [12–17]. Quality-aware multimodal fusion (QMF) provides a published failure-aware comparator by deriving modality confidence from predictive energy and aligning confidence with historical loss trajectories [18]. These approaches establish useful design patterns, but they do not make an arbitrary weighting response physically calibrated or comparable across sensors.

Sensor-fault diagnosis has a complementary engineering literature. Redundant sensor networks can localize offset, drift, noise, peaks and stuck responses, and can distinguish local sensor faults from global system changes [19–21]. Hydraulic condition-monitoring studies further show the value of heterogeneous pressure, flow, temperature, power and vibration measurements for component-state classification [22–25]. These studies motivate physically defined sensor groups and independent condition-monitoring validation rather than arbitrary computational groupings alone.

Most open benchmarks nevertheless encode sensor failure through simulation, laboratory intervention or post hoc feature corruption. A recent public AHU resource instead provides one year of building-management-system records from an office, an auditorium and a hospital, with return- and supply-air temperature sensor faults annotated by experienced HVAC engineers [26, 27]. These labels create a useful external field test, although they are rule- and expert-adjudicated operational annotations rather than maintenance-confirmed hardware replacement events.

Recent AHU studies show that the direct diagnosis problem differs from fault-tolerant recovery. Transformer models can exploit long temporal dependencies within a continuously operating hospital, while attention-based transfer and temporal-contextual learning explicitly address building shift and limited labels [28–30]. These architectures preserve fault signatures to identify the failure itself. In contrast, fault-tolerant classification seeks to suppress corrupted evidence while retaining another target, so superiority in one task need not transfer to the other.

We study paired-degradation responsive fusion (PDRF). “Responsive” denotes an operational change in fusion weight after a controlled sensor degradation; it does not imply physical health, calibrated risk or sensor variance. Each training observation is paired with a view in which one available sensor group is degraded. The model learns a bounded response, and directional losses encourage the affected sensor’s response to rise and its fusion weight to fall. The formulation does not invoke an identifiable causal counterfactual. Predictive calibration is evaluated separately because discrimination, uncertainty and calibration are different properties [31–37].

The study makes four contributions. First, it defines a strict recovery estimand that includes only controlled interventions that reach an available sensor group and separates prevention, retained correction and selection denominators. Second, it distinguishes a correctness-disagreement conditional score from the deployment estimand required on all observations. An all-row multinomial router, a two-stage disagreement-preference router, random routing, scalar confidence/entropy rules, a logit-margin rule and a shallow tree provide direct controls. Third, it separates lexicographic safety thresholds from thresholds that maximize a declared calibration-set utility. It also evaluates a post hoc proposed, deployment-level fitted-pair support gate retrospectively on the frozen experiment; the gate is intended for prospective use and mandates base-endpoint fallback when support is inadequate. Fourth, it evaluates selector transport across unseen mechanisms, calibration mixtures, corruption severities, affected sensor groups and highly imbalanced acquisition batches, while reporting fitting variation and engineering cost. The title and claims are limited to controlled corruption. Hydraulic results are single-rig sensitivity evidence, and expert-annotated air-handling-unit diagnosis remains a negative task boundary because it lacks an independent downstream recovery target.

2 Related work

2.1 Missing and unreliable sensor inputs

Missing-data mechanisms affect the interpretation of any robustness experiment [38]. Random sensor dropout tests tolerance to absent channels, whereas the present study targets channels that remain available but become corrupted. Corruption augmentation and consistency distillation can improve this setting, but they can also make an apparently novel fusion rule indistinguishable from robust training. Our comparator design therefore gives the main gates identical paired degradations, labels, minibatches and additional forward passes, and separately evaluates early and uniform fusion with clean distillation and modality-dropout self-distillation.

2.2 Bounded operational responses and probability calibration

Input-dependent losses can represent aleatoric variation under valid probabilistic assumptions, whereas PDRF uses its bounded response only for weighting [31, 32]. Response ordering and isotonic audits therefore do not convert it into physical variance. Probability quality is assessed separately with NLL, Brier score, multiple ECE definitions and temperature scaling because calibration depends on both multiclass structure and binning [37, 39, 40]. Entropy weighting, deep ensembles and conformal sets provide uncertainty-oriented comparators, not device-health estimates [33, 41]. Cross-building AHU controls likewise distinguish source-only GroupDRO from target-feature adaptation by CORAL or target normalization [42–44].

2.3 Selective prediction and model fallback

Selective classification and reject-option learning trade coverage for lower predictive risk [45–47]; recent work additionally studies selection under distribution shift [48]. Learning-to-defer methods route difficult cases to a human or another expert whose error pattern may differ from that of the primary predictor [49, 50]. Classifier cascades explicitly trade predictive accuracy against evaluation cost, while dynamic classifier and ensemble selection estimate which member is locally competent [51, 52]. Mixture-of-experts routing provides a related differentiable allocation mechanism [14]. The present fallback problem is narrower: both endpoints are fixed models of the same sensor-fusion task, no observation is rejected, and the selector specifically controls recovery-induced negative transfer while retaining corrections. Its operating curve is therefore prevention versus retained recovery, rather than conventional risk versus coverage, human deferral utility or average cascade cost.

2.4 Evidence and statistical estimands

Classifier comparisons can change with the resampling unit and evaluation estimand [53, 54]. Optimization seeds describe repeated fitting, whereas sample bootstrap intervals describe a fixed set of fitted predictors [55]. McNemar’s paired test evaluates discordant hard decisions on the same observations [56]. We report these levels separately and do not treat observations from one instrument as independent devices. Precision–recall metrics are included because class imbalance can make ROC summaries incomplete [57].

3 Methods

3.1 Problem definition

For observation i and sensor group m , let $\mathbf{x}_{im} \in \mathbb{R}^d$ be the feature vector, $a_{im} \in \{0, 1\}$ the availability indicator, $q_{im} \in (0, 1]$ an optional diagnostic-quality covariate and $y_i \in \{1, \dots, C\}$ the target. PDRF estimates

$$p(y_i \mid \{\mathbf{x}_{im}, a_{im}, q_{im}\}_{m=1}^M). \quad (1)$$

When no quality metadata are available, all q_{im} equal one. Every evaluated observation contains at least one available group. If all groups are unavailable, the implementation abstains before inference because normalized fusion is undefined.

For the routing audit, let b_i indicate that the base endpoint is correct, r_i that the recovery endpoint is correct and s_i that the router emits the recovery output. Table 1 fixes the event terminology used throughout.

Table 1: Hard-decision event definitions and denominators. N_C and N_H count correction and harm opportunities; n_C and n_H count retained corrections and realized harms. Prevention is $1 - n_H/N_H$, whereas recovery retention is n_C/N_C .

Quantity	Event definition	Denominator symbol	Reported rate
Correction opportunity	$\neg b_i \wedge r_i$	$N_C = \sum_i \mathbb{1}(\neg b_i \wedge r_i)$	n_C/N_C
Retained correction	$\neg b_i \wedge r_i \wedge s_i$	$n_C = \sum_i \mathbb{1}(\neg b_i \wedge r_i \wedge s_i)$	events / observations
Harm opportunity	$b_i \wedge \neg r_i$	$N_H = \sum_i \mathbb{1}(b_i \wedge \neg r_i)$	$1 - n_H/N_H$
Realized harm	$b_i \wedge \neg r_i \wedge s_i$	$n_H = \sum_i \mathbb{1}(b_i \wedge \neg r_i \wedge s_i)$	events / observations
Recovery endpoint selection	s_i	$N = \sum_i 1$	$\sum_i s_i/N$
Net correct change	$n_C - n_H$	N	events per 10,000

3.2 Architecture and bounded weighting

Each sensor group is encoded by affine layers $d \rightarrow 64 \rightarrow 32$, with a ReLU after each hidden affine transformation and no normalization or dropout. A biased linear $32 \rightarrow C$ output produces group-specific logits μ_{im} . The response head is exactly Linear(32, 16, bias = true), ReLU, then Linear(16, 1, bias = true). There is no activation after the final linear layer, so its raw response u_{im} can be positive or negative. The bounded operational response is

$$\begin{aligned} r_{im} &= B \tanh(u_{im}/B), \\ -B &< r_{im} < B. \end{aligned} \quad (2)$$

where the prespecified main setting uses $B = 3$. This layer placement matches the released implementation and gives the symmetric open interval in Eq. (2). Across 17,650 clean and faulted seed-observation responses from 10 fitted RO-PDRF-Full models, the observed float32 range was -2.899454 to 3.000000 ; the latter is numerical rounding of a value below the mathematical open boundary. Clean responses ranged from -2.757962 to 2.997449 . Unnormalized and normalized weights are

$$\begin{aligned} \tilde{w}_{im} &= a_{im} q_{im} \exp(-\beta r_{im}), \\ w_{im} &= \frac{\tilde{w}_{im}}{\sum_k \tilde{w}_{ik}}. \end{aligned} \quad (3)$$

with $\beta = 0.25$. Fused logits are $\ell_i = \sum_m w_{im} \mu_{im}$. The implementation applies a 10^{-6} lower guard to the normalization denominator, but the guard is inactive under the stated input constraints because every observation has at least one available group and $q_{im} \geq 0.01$ after clipping.

This construction differs from an unconstrained softmax gate in three testable ways. First, the weight is factorized into availability, an optional externally measured quality term and one bounded operational response; it is not an unrestricted function of the concatenated sensor representation. Second, the perturbed group is known during paired degradation, so its response and normalized weight receive directional supervision rather than only a task-gradient signal. Third, the same response that controls fusion is regularized by a group-specific predictive loss. Thus the response is operationally coupled to group evidence, although it is not asserted to be a physical health variable.

The factorization also gives two elementary stability properties. For an available group, fixed q and any $k \neq m$,

$$\begin{aligned} \frac{\partial w_{im}}{\partial r_{im}} &= -\beta w_{im}(1 - w_{im}) \leq 0, \\ \frac{\partial w_{ik}}{\partial r_{im}} &= \beta w_{ik} w_{im} \geq 0. \end{aligned} \tag{4}$$

For two available groups with equal quality, bounded responses imply

$$e^{-2\beta B} < \frac{\tilde{w}_{im}}{\tilde{w}_{ik}} < e^{2\beta B}. \tag{5}$$

At $B = 3$ and $\beta = 0.25$, one response alone can therefore create at most a 4.48-fold unnormalized-weight ratio between equal-quality groups. These properties do not guarantee that the learned ordering is correct; the experiments separately test that question.

The response regularizer, retained as $\mathcal{L}_{\text{score}}$ in the implementation for reproducibility, combines a bounded heteroscedastic-style term and a boundary penalty. For a minibatch $\mathcal{B}$, write $s_{im} = \omega_{y_i} [-\log \text{softmax}(\mu_{im})_{y_i}]$, where $\omega_c = n/(Cn_c)$ is the training-set class weight, and let $h_{im} = e^{-r_{im}} s_{im} + r_{im}$. Then

$$\mathcal{L}_{\text{score}} = \frac{\sum_{i \in \mathcal{B}} \sum_m a_{im} h_{im}}{\sum_{i \in \mathcal{B}} \sum_m a_{im}}, \tag{6}$$

$$\mathcal{L}_{\text{bound}} = \frac{\sum_{i \in \mathcal{B}} \sum_m a_{im} (r_{im}/B)^4}{\sum_{i \in \mathcal{B}} \sum_m a_{im}}. \tag{7}$$

Because cross-entropy is non-negative and $r_{im} > -B$, $\mathcal{L}_{\text{score}} > -B$. The complete objective therefore has a finite lower bound. This statement concerns optimization stability, not calibration of the bounded operational response.

3.3 Paired degradation and directional losses

For each minibatch observation, one available group m^* is drawn from a fixed training prior. A retention factor $\rho_i \sim \text{Uniform}(0.15, 0.50)$ controls additive Gaussian degradation with scale $2.5(1 - \rho_i)$. In half of the pairs, q_{im^*} is multiplied by ρ_i ; otherwise it is unchanged. Let primed symbols denote the degraded view and $d_i = 1 - \rho_i$.

The clean and degraded task losses use the class-weighted mean implemented by PyTorch. With $e_i(\ell) = -\log \text{softmax}(\ell)_{y_i}$,

$$\begin{aligned} \mathcal{L}_{\text{task}} &= \frac{\sum_{i \in \mathcal{B}} \omega_{y_i} e_i(\ell_i)}{\sum_{i \in \mathcal{B}} \omega_{y_i}}, \\ \mathcal{L}_{\text{deg}} &= \frac{\sum_{i \in \mathcal{B}} \omega_{y_i} e_i(\ell'_i)}{\sum_{i \in \mathcal{B}} \omega_{y_i}}. \end{aligned} \tag{8}$$

173 Let $\mathcal{V}_{\mathcal{B}} = \{i \in \mathcal{B} : \sum_m a_{im} > 1\}$ denote observations for which the selected group can be re-
 174 moved while another group remains. For $i \in \mathcal{V}_{\mathcal{B}}$, the removal prediction $\mathbf{p}_i^{(-m^*)}$ defines a severity-
 175 interpolated target $\mathbf{t}_i$; write $\hat{\mathbf{t}}_i = \text{stopgrad}(\mathbf{t}_i)$:

$$\begin{aligned} \mathbf{t}_i &= (1 - d_i)\mathbf{p}_i + d_i\mathbf{p}_i^{(-m^*)}, \\ \mathcal{L}_{\text{cons}} &= \frac{1}{|\mathcal{V}_{\mathcal{B}}|} \sum_{i \in \mathcal{V}_{\mathcal{B}}} \|\mathbf{p}'_i - \hat{\mathbf{t}}_i\|_2^2. \end{aligned} \tag{9}$$

176 If $\mathcal{V}_{\mathcal{B}}$ is empty, $\mathcal{L}_{\text{cons}}$ and the removal-dependent directional terms are omitted, equivalently set
 177 to zero. Directional terms operate on the affected normalized weight and raw response and are
 178 averaged over eligible observations. With $[z]_+ = \max(0, z)$ and $\Delta w_i = w'_{im^*} - w_{im^*}$,

$$\mathcal{L}_{\text{mono}} = |\mathcal{V}_{\mathcal{B}}|^{-1} \sum_{i \in \mathcal{V}_{\mathcal{B}}} [\Delta w_i]_+. \tag{10}$$

179 With $\gamma_i = 0.5d_i/0.85$ and $\Delta r_i = \text{stopgrad}(r_{im^*}) - r'_{im^*}$, the response-ranking term is

$$\mathcal{L}_{\text{rank}} = |\mathcal{V}_{\mathcal{B}}|^{-1} \sum_{i \in \mathcal{V}_{\mathcal{B}}} [\gamma_i + \Delta r_i]_+. \tag{11}$$

180 The final loss is

$$\begin{aligned} \mathcal{L} &= \mathcal{L}_{\text{task}} + 0.20\mathcal{L}_{\text{score}} \\ &\quad + 0.50\mathcal{L}_{\text{bound}} + 0.15\mathcal{L}_{\text{deg}} \\ &\quad + 0.10\mathcal{L}_{\text{cons}} + 0.05\mathcal{L}_{\text{mono}} \\ &\quad + 0.15\mathcal{L}_{\text{rank}}. \end{aligned} \tag{12}$$




Figure 1: **Bounded responsive fusion and recovery supervision.** **a**, Each sensor group produces class logits and a bounded operational response. Availability, optional diagnostic quality and the response determine normalized fusion weights. Distinct connector ports separate the logit and weight contributions to the fused prediction. **b**, Clean and faulted views share model parameters. The detached clean prediction supervises recovery of the faulted prediction; proper-scoring and ranking terms are auxiliary controls.

3.4 Comparators and mechanism ablations

Early fusion concatenates zero-filled group features, masks and quality values before a $128 \rightarrow 64 \rightarrow C$ multilayer perceptron; uniform fusion averages group logits. Both receive the same paired degradation and degraded-task loss as PDRF. CAGF concatenates four 32-dimensional group representations, masks and quality values before an unrestricted $64 \rightarrow M$ softmax gate. It shares PDRF’s paired-degradation, consistency and normalized-weight monotonicity losses. Direct weight ranking adds a gate-weight margin, while the quality-only control uses $a_{im}q_{im}$ without a learned response.

RO-AT-GATE tests stronger cross-group interaction. It adds learned group identities to 32-dimensional tokens and applies one four-head attention layer with a 64-unit feed-forward block; unavailable groups are masked. A shared $32 + 2 \rightarrow 16 \rightarrow 1$ head maps each contextualized token, availability and quality to an unrestricted softmax gate. The 26,809-parameter model receives the same paired views and recovery objective as RO-CAGF. It is a sensor-group attention control, not a reproduction of a task-specific vision or language system [6, 12, 13].

RO-MER supplies a sample-dependent multi-expert control motivated by failure-resilient routing [14]. It combines one classifier per physical group with a joint 64-unit expert and a separate 64-unit router. The prediction uses the resulting $(M + 1)$ -expert distribution. Losses requiring physical-group weights use the differentiable allocation

$$w_{im}^{\text{eff}} = \alpha_{im} + \alpha_{iJ} \frac{a_{im}}{\sum_k a_{ik}}. \quad (13)$$

The joint probability is distributed uniformly only for shared monotonicity and gate regularization; Eq. (13) does not alter the prediction. RO-MER has 35,811 parameters and uses the same faults, batches, seeds and checkpoint rule as RO-CAGF and RO-PDRF-Full.

The complete RO-MER objective is

$$\begin{aligned}\mathcal{L}_{\text{RO-MER}} = & \mathcal{L}_{\text{task}} + 0.15\mathcal{L}_{\text{deg}} + 0.10\mathcal{L}_{\text{cons}} + 0.05\mathcal{L}_{\text{mono}} \\ & + 0.30\mathcal{L}_{\text{rec}} + 0.05\mathcal{L}_{\text{Brier}} + 0.10\mathcal{L}_{\text{AUC}} + 0.005\mathcal{L}_{\text{router}},\end{aligned}\tag{14}$$

Here $\mathcal{L}_{\text{router}}$ is the clean-plus-degraded Kullback–Leibler divergence from $\mathbf{w}_i^{\text{eff}}$ to uniform allocation over available groups. RO-CAGF uses the same coefficients with native softmax weights; RO-PDRF-Full substitutes its response-interior penalty. Recovery supervision, proper scoring, pairwise ranking, paired views and task losses are otherwise matched.

We additionally adapt Quality-aware Multimodal Fusion (QMF) to the same four sensor groups while retaining its detached energy-confidence fusion and historical-loss ranking rule [18]. This QMF-PD control shares the present encoders, corruptions, optimizer, split and checkpoint criterion; it is a sensor-group adaptation rather than an exact reproduction of the original backbones or tasks. Implementation choices and sensitivity variants are specified in the Supplementary Information.

PDRF contains 19,740 trainable parameters on the chemical array, compared with 26,588 for CAGF and DWR, 35,811 for RO-MER, 26,182 for early fusion and 17,560 for uniform fusion. Ablations comprise paired degradation only (PD-ONLY), paired degradation plus normalized-weight monotonicity (PD-MONO), paired degradation plus response ranking (PD-RANK), the complete objective without $\mathcal{L}_{\text{score}}$ (NO-LSCORE), and $\mathcal{L}_{\text{score}}$ without $\mathcal{L}_{\text{rank}}$ (LSCORE-NO-LRANK). A PDRF-NOQ model is trained and tested with all quality values fixed to one.

3.5 Recovery-oriented extension

Reducing an affected group’s weight did not ensure prediction recovery. RO-PDRF therefore retains the PDRF architecture and distills the degraded prediction from the detached clean-view prediction. Writing $\hat{\mathbf{p}}_i = \text{stopgrad}(\mathbf{p}_i)$ gives

$$\mathcal{L}_{\text{rec}} = \frac{1}{|\mathcal{V}_{\mathcal{B}}|} \sum_{i \in \mathcal{V}_{\mathcal{B}}} D_{\text{KL}}[\hat{\mathbf{p}}_i \parallel \mathbf{p}'_i].\tag{15}$$

$\mathcal{L}_{\text{rec}}$ requires at least two available groups and is omitted when $\mathcal{V}_{\mathcal{B}}$ is empty. This excluded $0.386\% \pm 0.052\%$ of 5,933 training rows across seeds. PDRF plus $0.30\mathcal{L}_{\text{rec}}$ defines RO-PDRF-Lite. Label-assisted, confidence-weighted, exponential-moving-average and agreement-gated teachers remain supplementary controls; the deployed selector uses no held-out label [58].

Second, multiclass Brier losses on clean and degraded predictions provide proper-scoring supervision,

$$\begin{aligned}\mathcal{L}_{\text{Brier}} = & \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \|\mathbf{p}_i - \mathbf{e}_{y_i}\|_2^2 \\ & + \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \|\mathbf{p}'_i - \mathbf{e}_{y_i}\|_2^2.\end{aligned}\tag{16}$$

The Brier term is unweighted; class weights remain confined to cross-entropy. A weak log barrier then discourages boundary occupancy without removing the quartic penalty. With $\epsilon = 10^{-4}$, define $b_{\epsilon}(r) = -\log\{\max[\epsilon, 1 - (r/B)^2]\}$, $\bar{b}_{im} = b_{\epsilon}(r_{im}) + b_{\epsilon}(r'_{im})$ and $Z_{\mathcal{B}} = \sum_{i \in \mathcal{B}} \sum_m a_{im}$:

$$\mathcal{L}_{\text{interior}} = Z_{\mathcal{B}}^{-1} \sum_{i \in \mathcal{B}} \sum_m a_{im} \bar{b}_{im}.\tag{17}$$

Unavailable groups are excluded by a_{im} ; the clamp is numerical. Finally, a smooth one-versus-rest pairwise loss supervises degraded-view class ordering. Let P_c and N_c denote positive and negative minibatch observations, $\mathcal{C}_B = \{c : |P_c| > 0, |N_c| > 0\}$ and $\phi_{ijc} = \text{softplus}[-(\ell'_{ic} - \ell'_{jc})]$:

$$\mathcal{L}_{\text{AUC}} = \frac{1}{|\mathcal{C}_B|} \sum_{c \in \mathcal{C}_B} A_c, \quad (18)$$

$$A_c = \frac{\sum_{i \in P_c} \sum_{j \in N_c} \phi_{ijc}}{|P_c| |N_c|}.$$

The pairwise term is omitted when no class has both outcomes [59]. RO-PDRF-Full adds $0.30\mathcal{L}_{\text{rec}} + 0.05\mathcal{L}_{\text{Brier}} + 0.005\mathcal{L}_{\text{interior}} + 0.10\mathcal{L}_{\text{AUC}}$ to Eq. (12). RO-PDRF-Lite is the practical default, Full is the mechanistic architecture comparator and RO-PDRF-EMA is a Full-based calibration sensitivity. No test observation fits a model or ensemble parameter.

To separate recovery supervision from architecture, we also train recovery-oriented CAGF (RO-CAGF). It receives the same clean-view distillation, Brier and degraded-view ranking losses as RO-PDRF-Full. Its unrestricted softmax gate uses a weak KL-to-uniform regularizer (weight 0.005), whereas RO-PDRF-Full uses the response-interior penalty. The matched data, fixed faults, seeds, minibatches and checkpoint rules isolate architecture subject to this necessary regularizer difference. Regularizer grids, early/uniform-fusion controls, modality-dropout self-distillation and confidence-thresholded distillation are reported in the Supplementary Information [5, 6].

Endpoint identities were frozen before selector test evaluation. PDRF is the stable base because it contains no recovery-distillation pathway and requires one ordinary forward pass. RO-PDRF-Lite is the primary recovery endpoint because it adds one explicit recovery term without changing parameter count or inference architecture. RO-PDRF-Full improved affected-subset AUROC by only 0.002 over Lite in the matched endpoint benchmark and is retained as a mechanistic loss-bundle audit rather than silently replacing the practical endpoint. Table 2 places this choice beside architecture-matched and published-method controls; the table documents the frozen design and was not used to tune the selector on test outcomes.

Table 2: **Concise endpoint benchmark and routing role.** Affected-subset macro-AUROC values summarize the matched chemical-array endpoint study over 10 optimization seeds. Parameter counts exclude the downstream logistic selector. Endpoint identity was frozen before the selective-recovery test audit.

Method	Role in endpoint benchmark	Parameters	Affected AUROC	Router endpoint
PDRF	stable base reference	19,740	0.791	Base
RO-PDRF-Lite	practical paired-recovery model	19,740	0.802	Recovery
RO-PDRF-Full	mechanistic recovery audit	19,740	0.804	No
RO-CAGF	architecture-matched learned baseline	26,588	0.797	No
RO-AT-GATE	attention-gated learned baseline	26,809	0.784	No
QMF-PD	published-method adaptation	17,560	0.775	No

3.6 Cross-fitted selective recovery

Unconditional recovery can repair a wrong base prediction but can also overturn a correct one. We therefore treat routing as hard-decision risk control. Here “Lite” means that the router selects between PDRF and the RO-PDRF-Lite recovery endpoint; it does not merely denote a smaller selector. Both endpoints require one ordinary forward pass and have the same 19,740-parameter architecture. The heavier PDRF/RO-PDRF-Full pair is retained only as a supplementary mechanistic and legacy-routing sensitivity.

Batch 7 is divided chronologically into two disjoint halves. The first half fits one scalar temperature per endpoint by minimizing multiclass negative log-likelihood. The second contains two calibration-only realizations of each fault mechanism; only interventions that reach an available affected group enter selector development. Routing features use the temperature-scaled endpoint probabilities. The selected endpoint vector is emitted without probability mixing. Uncalibrated endpoint features, Platt scaling and isotonic calibration of the preference score are separate controls rather than replacements selected from test performance.

Lite-CF uses six prespecified inputs: base confidence, recovery confidence, their difference, base normalized entropy, recovery normalized entropy and their difference. It receives neither the class label nor fault type at inference. Let z_{ij} denote standardized feature j . Feature means and standard deviations are fitted only on the informative training rows inside each cross-fitting fold, then refitted on all informative calibration rows for the final model. An L_2 -regularized logistic model produces the *conditional preference score*

$$\pi_i = \sigma \left(\gamma_0 + \sum_{j=1}^6 \gamma_j z_{ij} \right), \quad (19)$$

using calibration rows on which PDRF and RO-PDRF-Lite differ in correctness. The target equals one when recovery is correct and the base endpoint is wrong. Therefore π_i estimates a conditional preference on the label-selected subset $b_i \neq r_i$; it does not estimate $\Pr(r_i = 1, b_i = 0 \mid \mathbf{x}_i)$ or a general reliability probability for all deployment rows. Lite-CF uses $C = 0.05$, the `lbfgs` solver, balanced class weights, tolerance 10^{-4} , 2,000 maximum iterations and random state 0. We retain this model as the frozen historical router and explicitly compare it with all-row alternatives.

Coefficient fitting is separated from threshold selection by five-fold stratified group cross-fitting. Realizations and mechanisms from the same calibration observation remain in one fold. Candidate thresholds 0.50, 0.525, $\dots$, 0.90 are evaluated only on pooled out-of-fold scores. Selection minimizes realized harms, then maximizes selected accuracy and retained corrections; remaining ties use the highest threshold. The formal fits had a mean of 1.5 tied thresholds (range 1–4). Final coefficients are refitted on all informative rows while the threshold remains frozen.

For future prospective deployment, mathematical fit is necessary but insufficient. We subsequently specified a conservative reference support gate requiring at least 25 events in each preference class, median out-of-fold AUROC at least 0.65 across 10 repeated grouped splits, no repeated AUROC below 0.60 and a threshold range no wider than 0.10. Both preference classes must also occur in every fitting fold. These support criteria were proposed after the frozen experiment and applied retrospectively to the frozen fits; they were not prespecified and do not replace the ungated 94.7%/9.7% estimate. The retrospective analysis defines the intended prospective rule: a site that fails any condition emits PDRF while collecting new calibration evidence. Sensitivity varies the minimum event count over 15, 25 and 40, median AUROC over 0.60, 0.65 and 0.70, and maximum threshold range over 0.05, 0.10 and 0.15, while holding the every-fold class requirement and minimum repeat AUROC of 0.60 fixed. The released routing API returns probabilities by default and can optionally return selected endpoint identity and the conditional preference score; routed

confidence remains endpoint-specific rather than globally calibrated.

We address the selected-sample problem with an all-row outcome $o_i \in \{0, 1, 2\}$ denoting base better, equivalent hard-decision correctness or recovery better. A multinomial logistic model produces probabilities $\eta_{ik} = \Pr(o_i = k \mid \mathbf{x}_i)$ and score $d_i = \eta_{i2} - \eta_{i0}$. A two-stage alternative first estimates $\Pr(b_i \neq r_i \mid \mathbf{x}_i)$ on all rows and then multiplies it by the conditional base/recovery preference. Both stages are cross-fitted by original observation. These models target the unconditional decision problem without inverse-probability correction because all calibration rows and their three outcomes are observed.

Controls use the same two endpoint passes and calibration budget. They include random ranking, low base confidence, recovery-minus-base confidence, base-minus-recovery entropy, recovery-minus-base top-two logit margin, a depth-two tree with minimum leaf size 35, rank-normalized logistic features and class-conditional endpoint-prediction indicators. Complete operating curves rank scores separately within each fixed test mechanism-fitted-pair cell and use the corresponding test labels to calculate prevention and retention. They are descriptive cell-stratified test ranking audits, not deployable operating policies. For matched controls, thresholds are selected only on calibration to match the frozen router’s calibration prevention or retention and are then applied unchanged to test rows.

We further probe sparse estimation through ten repeated grouped splits for each endpoint pair, standardized coefficient/sign distributions, $C \in \{0.01, 0.05, 0.1, 0.5, 1, 5\}$, calibration-size subsampling and five calibration-family mixtures (uniform and one 70%-dominant mechanism). Leave-one-mechanism-out and Gaussian-only calibration test mechanism transport. Severity transport fits selectors at scales 1, 2 or 3 and evaluates every scale without endpoint refitting; affected-group transport calibrates on group 1 and evaluates groups 1–4 at scale 3.

The deployment audit reuses the same endpoints, temperatures, selector and test realization for every stream. It includes the strict fault-applied-and-available subset, a complete 40%-prevalence mixed stream and a no-imposed-fault stream with the same fixed 20% group-missingness process. Probability quality is reported separately for endpoint-class agreement and disagreement. The primary outcomes are negative-transfer prevention, recovery retention and

$$U = V_C n_C - C_H n_H - \lambda(P - 1), \quad (20)$$

where n_C and n_H are retained corrections and realized harms per 10,000 observations, V_C is correction value, C_H is harm cost, P is mean forward-pass count and λ is a pass-cost proxy in correction-equivalents per 10,000 observations. We vary C_H/V_C over 1, 2, 5 and 10, imposed-fault prevalence and λ . Energy was not measured; forward passes and counted linear-layer FLOPs remain computational proxies.

Two threshold regimes are kept distinct. The safety threshold uses the frozen lexicographic rule: minimize calibration harms, then maximize selected accuracy and retained corrections. The application threshold maximizes $n_C - (C_H/V_C)n_H$ on pooled out-of-fold calibration predictions for each declared cost ratio. Both are evaluated once on test. Always-base and always-recovery remain explicit endpoints. Reporting includes extra forward passes per retained and net correction, selected observations per net correction, model storage, router latency and orchestration residual beyond the sum of endpoint timings.

The deployment hierarchy is explicit. First, the fitted PDRF/RO-PDRF-Lite pair must pass the local pair-level support gate. Second, a supported pair may use conditional Lite-CF with its calibration-frozen lexicographic threshold; an unsupported pair falls directly back to PDRF and cannot switch automatically to another router family. Third, all-row and two-stage utility thresholds are calibration-frozen analytical controls indexed by a declared harm ratio. They do not have

independently validated model-specific support gates in this study and are therefore not claimed as supported deployment policies. Prospective use of either would require a separately frozen support rule; failure of that rule would also revert to PDRF. Consequently, the observed two-stage advantage at a two-to-one ratio is a fixed-test analytical result, not a deployment recommendation.

3.7 Chemical-sensor dataset and fixed faults

The UCI Gas Sensor Array Drift Dataset contains 13,910 observations from 16 metal-oxide sensors, six gas classes and 10 acquisition batches over 36 months [60, 61]. Each sensor contributes eight extracted response features. The principal grouping uses four consecutive sensors per group. All models receive the same grouping. Sensitivity analyses additionally use interleaved and five fixed random partitions; no partition is called a physical modality.

Batches 1–6 train the models. Their feature means and standard deviations are the only quantities used for standardization. The same transform is applied to batches 7–10, after which each feature is clipped to $[-8, 8]$ before sensor grouping. Thus every reported “training range” is the empirical minimum–maximum after training-only standardization and clipping; it is not a physical sensor limit. The first 60% of Batch 7 selects checkpoints, and the remaining 40% is reserved for calibration; the routing experiment further divides that calibration partition into temperature and selector halves. Batches 8–10 are never used to fit preprocessing, endpoints, temperatures, selectors or thresholds.

Controlled faults are imposed after this training-only preprocessing in grouped standardized space. The formal test assigns 40% of observations a scale-3 intervention on group 1 and fixes $q = 1$. Fault assignment and 20% group missingness are generated independently to isolate their factorial effects. This is a stress design, not a claim that all missingness–corruption combinations occur naturally. Of 4,364 test observations, 1,765 are assigned the intervention; group 1 remains available and receives it in 1,436, whereas it is masked in 329. The strict recovery estimand is the former *fault-applied-and-available* subset. Fault indices, missingness masks and intervention values are fixed across methods and optimization seeds.

The complete test set contains 294, 470 and 3,600 observations in batches 8, 9 and 10. Its six class counts are 691, 685, 740, 708, 821 and 719. The corresponding strict-subset batch counts are 91, 165 and 1,180, and its Gaussian class counts are 230, 220, 230, 253, 266 and 237. Complete mechanism-by-batch class tables are supplied in the Supplementary Information.

Checkpoint selection and temperature calibration use separate fault-free views with independently generated 10% group missingness; no controlled Gaussian, offset, drift or stuck-at intervention is added to either subset. The checkpoint criterion is unweighted cross-entropy on the complete selection subset. For the cross-fitted deployment run, temperature is fitted on the first chronological half of the calibration partition, not an affected subset, by minimizing unweighted multiclass negative log-likelihood. The resulting scalar is specific to each fitted seed and method and is applied unchanged to every deployment stream. Controlled-fault and no-fault deployment streams use the same fixed 20% group-missingness process so that fault prevalence is the changing factor. Earlier natural-view analyses without imposed test missingness remain supplementary sensitivity evidence.

The expanded fault study crosses four prespecified mechanisms with realizations 70001–70008 and fitted seeds 101–105. Model fitting always uses the Gaussian paired degradation defined above; offset, drift and stuck-at are test-time interventions only. Gaussian faults add independent standardized noise. Offset faults add a featurewise signed displacement. Drift multiplies that displacement by a monotone 0.25–1.00 trajectory over held-out acquisition order. Stuck-at faults replace the affected group’s standardized features by a fixed signed vector. Within each realization, all

mechanisms share affected rows and missingness masks, while intervention values remain fixed across methods. These are controlled repeats on one instrument.

Table 3 translates the formal scale into empirical training-distribution displacement. At scale 3, median absolute shifts were 1.89–3.10 training standard deviations and median percentile shifts were 0.36–0.56. Only 9.4–45.7% of corrupted values remained inside the training 1st–99th percentile interval. The formal setting is therefore a severe standardized stress test, not a typical natural-failure model. Scale-1 and scale-5 summaries and representative inverse-transformed feature examples are reported in the Supplementary Information.

Table 3: Fault-plausibility audit in the strict subset. Fractions are calculated featurewise against batches 1–6. “Training range” denotes the observed training minimum–maximum, not a physical sensor limit.

Mechanism (scale 3)	Median $ \Delta z $	Median percentile shift	Within training 1–99%	Within training range
Gaussian	2.015	0.360	0.379	0.609
Offset	3.000	0.538	0.332	0.632
Drift	1.888	0.440	0.457	0.697
Stuck-at	3.105	0.557	0.094	0.656

Exploratory controls vary diagnostic-quality metadata, leave-one-group-out training, response bounds, weighting slope, ranking weight, fault group, severity and prevalence. They share observations and fitted models, are not treated as independent replications and do not replace the prespecified main setting. Full specifications and results are reported in the Supplementary Information.

3.8 Independent hydraulic single-rig sensitivity analysis

The UCI Condition Monitoring of Hydraulic Systems dataset contains 2,205 raw 60-s cycles from an experimental hydraulic rig and native condition labels for cooler, valve, pump and accumulator states (DOI: 10.24432/C5CW21) [22, 62]. We retain 1,449 cycles marked as stable. Fourteen physical sensors form five groups by measured quantity: six pressure sensors, two flow sensors, four temperature sensors, motor power and vibration. This grouping is defined before model fitting and follows the measurement system rather than a random partition.

Each sensor is summarized by distributional and trend statistics. Unequal group lengths are aligned to 48 inputs; a zero-padded native-order representation audits possible smoothing. The main split preserves acquisition order within every native condition class, assigns contiguous 60%/15%/10%/15% blocks and removes five cycles at each boundary. A global chronological split would omit some condition classes from training. After guard-band removal, the test counts are 72/72/74 for the three cooler states, 54/54/54/56 for the four valve states, 74/72/72 for the three pump states and 56/54/54/54 for the four accumulator states. Recovery-objective-matched RO-PDRF-Full and RO-CAGF use five matched seeds; controlled pressure- and vibration-group faults affect the same 40% of test cycles. Component-condition labels are native to the rig, whereas sensor faults remain interventions. Representation and random-split sensitivities are supplementary [23–25].

3.9 Real-operational AHU direct-diagnosis boundary analysis

We additionally use the open real-operational AHU dataset from an auditorium, a hospital and a large office (Figshare DOI: 10.6084/m9.figshare.27147678.v3) [26, 27]. The binary task uses

native, expert-annotated temperature-sensor-fault labels and no artificial test corruption. Seven measurements form five physical groups; each building uses a chronological 60%/15%/10%/15% split and five matched seeds. The untouched test periods contain 228/9,354 fault records in the auditorium (2.44%), 1,495/9,362 in the hospital (15.97%) and 346/19,115 in the office (1.81%). Return-air/supply-air subtype counts are 179/49, 1,487/8 and 293/53, respectively. Because the label is the sensor fault itself and no independent downstream state is available, this is a negative task-boundary audit, not downstream recovery validation. Detailed shift, adaptation, risk-coverage and conformal analyses are reported in the Supplementary Information [41–44].

3.10 Training, metrics and statistical analysis

AdamW uses learning rate 2×10^{-3} , weight decay 10^{-4} , batch size 256 and gradient clipping at 5. Chemical-array training uses at most 60 epochs, hydraulic training uses at most 80 and AHU training uses at most 12. Checkpoints minimize unweighted selection cross-entropy and stop after nine epochs without an improvement greater than 10^{-4} . RO-PDRF-Lite, RO-PDRF-Full, RO-PDRF-EMA and all teacher controls use this identical criterion; no test metric selects a checkpoint or variant. Class weights are $n/(Cn_c)$. No dropout or batch normalization is used.

The primary deployment outcomes are negative-transfer prevention, recovery retention and the utility in Eq. (20). Macro-AUROC, macro-AUPRC, NLL, multiclass Brier score and calibration are secondary. Equal-width ECE partitions maximum class probability into 15 intervals on $[0, 1]$. Equal-mass ECE splits the sorted maximum probabilities into 15 groups. Classwise ECE applies 15 equal-width bins to each one-versus-rest probability and averages the six values. Classwise AUROC, AUPRC and recall are archived. Temperature scaling is fitted separately for each endpoint and optimization seed on the disjoint first half of Batch 7.

The deployment analyses are frozen as CEE-CF10-R2 for the Full endpoint and CEE-CF10-R2-LITE for the primary Lite endpoint. Within each seed, one endpoint pair is reused across all four faults and selector comparisons. Test realization 70001 is fixed. Means over the strict set give equal weight to the four mechanism-seed cells. Clean/no-fault summaries contain one value per fitted pair. Mechanism-specific effects use paired Student- t intervals across the 10 fitted model pairs while holding the mechanism fixed. They are labelled fitted-pair or fitting-variation intervals; all source files also report every point, median, interquartile range and full range. No interval is interpreted as a significance test.

The strict subset contains 91 Batch 8, 165 Batch 9 and 1,180 Batch 10 observations per mechanism, so Batch 10 supplies 82.2% of the observation-weighted result. The complete mechanism $\times$ batch $\times$ method matrix reports accuracy, macro-AUROC, correction and harm opportunities, prevention, retention and utility. We separately report pooled observation-weighted performance, an equal-batch average, the worst acquisition batch and pooled Batches 8–9 versus Batch 10. Batch-aware intervals hold mechanism, batch and test rows fixed and vary only the 10 fitted endpoint pairs.

The previous four-mechanism crossed bootstrap is retained only as a labelled sensitivity in the Supplementary Information. It is not used for new headline inference because four top-level mechanism clusters do not support a reliable mechanism-population bootstrap. Exact Wilcoxon tests across optimization seeds are likewise reported only as fitting-stability diagnostics. Hydraulic summaries use a task-fault-seed hierarchy; AHU results remain building-specific. The strict estimand and selector analyses followed reviewer-driven audit and are not described as preregistered confirmatory analyses.

Experimental protocol, fixed interventions and evidence boundary




Figure 2: **Integrated experimental protocol and controlled fault design.** The upper lane follows the chemical-array acquisition order from batches 1–6 training through Batch 7 checkpoint selection/calibration to held-out batches 8–10; fault indices, masks and intervention values are fixed across methods and seeds. The lower lane follows the independent hydraulic rig from stable 60-s cycles to the guarded within-condition blocked main analysis; the earlier stratified split is retained only as a sensitivity analysis. Component-condition labels are native to the rig, whereas chemical, pressure and vibration sensor corruptions remain controlled interventions.

4 Results

4.1 Always-base context and the retrospective support-gate analysis

The strict subset contains 1,436 observations per mechanism on which the assigned intervention reached an available group. Across the 40 fixed mechanism–fitted-pair cells, PDRF achieved accuracy 0.4869 and macro-AUROC 0.8115. Unconditional RO-PDRF-Lite increased these values to 0.5002 and 0.8175, but it realizes every possible correction and every possible harm. Always PDRF is therefore the maximum-safety reference: 100% negative-transfer prevention, 0% recovery retention and zero *incremental routing* utility relative to PDRF by definition (Table 4); this is not the total utility of the underlying classification system.

The historical frozen Lite-CF result is explicitly ungated. It selected recovery on 52.2% of strict rows, prevented 94.7% of recovery-induced negative transfers and retained 9.7% of available corrections, yielding accuracy 0.4885, macro-AUROC 0.8151 and equal-cost utility +15.8 per 10,000. Only 5/10 fitted endpoint pairs passed the subsequently specified support criteria applied retrospectively to the frozen fits. Retrospective fallback of unsupported pairs to PDRF reduced recovery selection to 31.1%, increased prevention to 96.2% and reduced retention to 6.6%. Its

macro-AUROC was 0.8132 and equal-cost utility was +7.5. This post hoc analysis is safer than the frozen router but only 0.0017 AUROC above always PDRF on this fixed test set; it illustrates the proposed future rule rather than constituting a prospective validation.

Table 4: **Strict-subset endpoints and calibration-frozen policies.** Values are means over 40 fixed mechanism-fitted-pair cells. Prevention and retention have different opportunity denominators. $U_1/10^4$ is incremental routing utility relative to PDRF with equal correction and harm values and no compute penalty. The support-gated row is a post hoc gate applied retrospectively to frozen fits; the 94.7%/9.7% row is the ungated frozen result. The two-stage 1:1 row is included for completeness but remains an analytical control.

Policy	Accuracy	Macro-AUROC	Prevent (%)	Retain (%)	Select (%)	$U_1/10^4$
Always PDRF	0.4869	0.8115	100.0	0.0	0.0	+0.0
Always RO-PDRF-Lite	0.5002	0.8175	0.0	100.0	100.0	+132.1
Frozen Lite-CF (ungated)	0.4885	0.8151	94.7	9.7	52.2	+15.8
Retrospective support-gated Lite-CF	0.4877	0.8132	96.2	6.6	31.1	+7.5
All-row multinomial (safety control)	0.4879	0.8118	97.2	7.3	3.0	+9.6
All-row multinomial (utility 1:1)	0.4977	0.8160	37.5	70.3	66.8	+107.9
Two-stage (utility 1:1)	0.4980	0.8165	29.4	75.4	80.2	+110.6

Across the full $3 \times 3 \times 3$ support-gate grid, 1–7 of 10 fitted pairs passed; mean prevention ranged from 95.2% to 97.9%, retention from 1.4% to 8.6% and equal-cost utility from +0.35 to +13.75 per 10,000. The conservative reference configuration admitted 5/10 pairs and had worst-batch utility −2.54, where each fitted pair first contributes its worst acquisition-batch mean. Tightening the event minimum to 40 admitted only one pair, whereas relaxing all three varied criteria admitted seven. Several one-factor changes left the admitted set or decision outcomes unchanged (Table 5), showing that 25/0.65/0.10 is a transparent reference, not a uniquely correct operating standard.

Table 5: **One-factor sensitivity of the post hoc deployment-support gate.** The reference row (bold passed count) requires 25 events per preference class, median repeated AUROC 0.65 and threshold range at most 0.10. Every-fold class support and minimum repeat AUROC 0.60 remain fixed. Worst-batch utility averages each fitted pair’s least favourable acquisition-batch mean. The complete 27-configuration grid is provided in the Supplementary Information.

Min. events	Median AUROC	Max. Δt	Passed	Select (%)	Prevent (%)	Retain (%)	$U_1/10^4$	Worst-batch $U_1/10^4$
15	0.65	0.10	6/10	36.7	96.2	6.6	+7.49	-2.54
25	0.60	0.10	5/10	31.1	96.2	6.6	+7.49	-2.54
25	0.65	0.05	3/10	21.6	97.0	3.4	+3.31	-3.17
25	0.65	0.10	5/10	31.1	96.2	6.6	+7.49	-2.54
25	0.65	0.15	5/10	31.1	96.2	6.6	+7.49	-2.54
25	0.70	0.10	5/10	31.1	96.2	6.6	+7.49	-2.54
40	0.65	0.10	1/10	7.8	97.9	1.4	+0.35	-0.42

The 52.2% selection rate and 9.7% retention rate have different denominators. Selection is normalized by all 1,436 strict observations, whereas retention is normalized only by N_C , the sparse rows on which recovery alone corrects PDRF. Selection can therefore be high because the router also chooses recovery on endpoint-class agreements and rows where both endpoints share correctness status. Table 6 reports the absolute event counts rather than only opportunity-normalized rates.

The routed probability vector improved on PDRF for NLL, Brier score, ECE and macro-AUPRC, but remained below unconditional RO-PDRF-Lite for every strict-set probability metric. On the complete 40%-prevalence stream, Lite-CF/PDRF/RO-PDRF-Lite macro-AUROC values were 0.8441/0.8426/0.8451, and NLL values were 2.1386/2.2131/2.1135. On the no-

Table 6: **Absolute strict-subset opportunities and routed events by mechanism.** N_C and N_H are correction and harm opportunities; n_C and n_H are retained corrections and realized harms. Counts are mean absolute events per 1,436-observation fitted pair, with ranges across 10 fitted pairs in parentheses. Select uses all observations; prevent and retain use N_H and N_C , respectively.

Mechanism	N_C	N_H	n_C	n_H	Select (%)	Prevent (%)	Retain (%)
Gaussian	45.9 (36–70)	18.9 (13–32)	6.2 (0–16)	1.5 (0–4)	51.9	91.2	13.5
Offset	33.6 (10–49)	16.9 (3–37)	2.7 (0–9)	1.9 (0–15)	54.1	93.6	8.1
Drift	24.7 (5–43)	13.6 (4–29)	1.9 (0–7)	0.7 (0–6)	53.9	97.3	11.1
Stuck-at	38.7 (16–86)	17.6 (2–36)	3.1 (0–15)	0.7 (0–4)	48.7	96.6	6.1

imposed-fault stream, their macro-AUROC were 0.8629/0.8631/0.8623 and NLL values were 1.8636/1.9002/1.8626. Lite-CF prevented 94.2% of no-fault negative transfers but retained only 1.7% of no-fault corrections, giving -5.3 net correct decisions per 10,000. A low residual-harm rate—the observation-level form of a harmful switch—therefore does not imply that the routed probability vector is uniformly better.

4.2 All-row routing and simple controls exposed the selected-sample limit

Each calibration fit contained a mean of 36.1 base-better, 1,771.2 equivalent and 53.6 recovery-better rows. The all-row multinomial model used all three outcomes rather than fitting only the 89.7 correctness disagreements. Its pooled out-of-fold confusion matrix classified 73.7% of base-better, 81.8% of equivalent and 56.3% of recovery-better rows correctly; mean out-of-fold accuracy, NLL and three-class Brier score were 0.809, 0.523 and 0.260. The model therefore reduced the estimand mismatch but did not make the rare endpoint-preference classes easy to identify.

At its safety threshold, the all-row router achieved 97.2% prevention and 7.3% retention with only 3.0% recovery selection and $+9.6$ equal-cost utility. Complete fixed-test curves additionally compare the conditional model, two-stage model, random ranking, confidence difference, entropy difference, base uncertainty, logit-margin difference and shallow tree (Fig. 3a). Under this descriptive cell-stratified test ranking audit, no tested score empirically dominated the frozen operating point in both mean prevention and mean retention. These curves are not deployable operating policies. In contrast, calibration-frozen matched thresholds are operationally reproducible comparisons: the random control at the calibration-prevention target produced 97.8% test prevention and 2.1% retention, whereas its calibration-retention target produced 79.4% prevention and 20.5% retention. The entropy rule gave 97.4%/7.5% and 85.3%/21.8% in the corresponding regimes. This transport error is why no simple comparison threshold replaces the frozen router.

Class-conditional logistic indicators improved pooled-row disagreement preference AUROC from 0.562 to 0.680 and produced $+85.0$ equal-cost utility, but prevention fell to 71.5%. A direct logit-margin utility rule produced $+118.4$ utility with only 19.2% prevention. Rank-normalized features retained 94.3% prevention but only 8.0% of corrections. These controls show that greater average utility is achievable by accepting more harmful transfers; they do not dominate the safety policy.

Mechanism-fixed paired intervals replace the four-cluster bootstrap (Table 7). On the present fixed test set, Gaussian accuracy, AUROC and equal-cost utility differences were positive across fitted model pairs. Stuck-at AUROC was also positive, whereas its accuracy and utility intervals crossed zero. Offset and drift intervals crossed zero for most hard-decision quantities. Gaussian AUROC differences were positive within batches 8, 9 and 10 ($+0.0062$, $+0.0077$ and $+0.0061$). These intervals quantify fitting variation on fixed observations, not a population of devices or fault mechanisms.

Table 7: **Mechanism-fixed paired differences for the hard-decision objective.** Values are Lite-CF minus PDRF. Brackets are paired 95% Student- t intervals across 10 fitted model pairs while holding mechanism and test observations fixed. Equal-cost utility uses $V_C = C_H = 1$ and $\lambda = 0$. These are conditional fitting-stability summaries, not population intervals.

Fixed mechanism	Δ accuracy	Δ macro-AUROC	Equal-cost $U/10^4$
Gaussian	+0.0033 [+0.0005, +0.0061]	+0.0064 [+0.0034, +0.0093]	+32.73 [+4.94, +60.52]
Offset	+0.0006 [-0.0012, +0.0024]	+0.0028 [0.0000, +0.0057]	+5.57 [-12.44, +23.58]
Drift	+0.0008 [-0.0006, +0.0022]	+0.0018 [-0.0002, +0.0037]	+8.36 [-5.69, +22.41]
Stuck-at	+0.0017 [-0.0004, +0.0037]	+0.0034 [+0.0005, +0.0063]	+16.71 [-3.92, +37.35]

The formal scale 3 intervention is a severe stress test, so we repeated the frozen test protocol at scale 1 without refitting endpoints, temperatures, selector coefficients or threshold. Lite-CF then reached 0.8357 macro-AUROC versus 0.8344 for PDRF, prevented 96.0% of negative transfers while retaining 5.3% of corrections and produced equal-cost utility +4.00 per 10,000 strict observations (Table 8). Milder corruption improved absolute endpoint performance but left fewer useful recovery opportunities, so safety prevention did not imply greater practical benefit.

Table 8: **Mild and formal controlled-corruption severity.** Values are equal-weight means over four fixed mechanisms and 10 fitted model pairs on the strict subset. The scale 1 audit reuses the scale 3-calibrated selector without refitting or retuning. $U/10^4$ uses equal correction and harm values with no pass penalty.

Scale	PDRF acc.	Lite-CF acc.	PDRF AUROC	Lite-CF AUROC	Prevent (%)	Retain (%)	$U/10^4$
1	0.5301	0.5305	0.8344	0.8357	96.0	5.3	+4.00
3	0.4869	0.4885	0.8115	0.8151	94.7	9.7	+15.84

4.3 Utility-aligned thresholds changed which policy was preferred

The frozen safety threshold and application-specific utility thresholds answer different questions. At equal correction and harm values, the frozen router yielded +15.8 and the retrospectively support-gated analysis +7.5 utility units per 10,000, whereas utility-selected all-row and two-stage analytical controls yielded mean test utilities of +107.9 and +110.6. Always recovery still gave the largest mean utility, +132.1, at the one-to-one ratio. When one harm cost two corrections, always recovery fell to +15.5, while calibration-selected all-row and two-stage routing yielded +30.8 and +37.8.

Fitted-pair dispersion materially qualifies the means. At the two-to-one ratio, the two-stage mean was +37.8 with a 95% fitted-pair interval of -2.7 to +78.3 and positive utility in 7/10 pairs. Its mean out-of-fold calibration utility was +136.8, giving a mean test-minus-calibration transport change of -99.0. All-row routing had mean +30.8 (-12.3 to +74.0; 7/10 positive), and conditional routing had +27.7 (-23.9 to +79.3; 7/10 positive). Thus, on this fixed test set and with $\lambda = 0$, two-stage had the highest mean utility among the tested policies at 2:1, but this is an analytical comparison rather than a supported deployment recommendation.

This advantage did not extend to high harm costs. At ratios 5 and 10, every learned or recovery policy had negative mean test utility and always PDRF, with zero incremental routing utility, dominated (Table 9; Fig. 3c). The all-row safety threshold was the least negative learned policy at ratio 5 (-8.5), but it still did not justify routing. Every application threshold was selected on pooled out-of-fold calibration predictions before test evaluation. The failure at high

ratios is therefore transport evidence, not a test-tuned threshold choice. Complete-stream utility also remained negative at 0% and 10% controlled-fault prevalence. Deployment must fix expected prevalence, V_C , C_H and compute cost before choosing a threshold.

Table 9: **Fitted-pair dispersion of utility-selected analytical controls.** Each row first pools the four fixed mechanisms within each of 10 fitted endpoint pairs. Utilities are units per 10,000 strict observations with $V_C = 1$, $\lambda = 0$ and the indicated harm ratio. Test values are mean [95% Student- t interval] and median [interquartile range]. “Test-cal.” compares fixed-test utility with utility on pooled out-of-fold calibration predictions. Positive counts denote fitted pairs with test utility above zero. Because decision-event denominators are pooled within fitted pair before averaging, decision percentages can differ slightly from the equal-weight mechanism-pair means in Table 4.

Ratio	Policy	Cal. U	Test U [95% CI]	Median [IQR]	Test-cal.	Positive	Select (%)	Prevent (%)	Retain (%)
1:1	Conditional	+180.5	+106.4 [+43.3,+169.5]	+106.2 [+33.1,+157.6]	-74.1	9/10	80.1	40.9	65.1
1:1	All-row	+194.8	+107.9 [+50.5,+165.4]	+107.1 [+38.7,+132.7]	-86.9	10/10	66.8	36.4	69.2
1:1	Two-stage	+183.7	+110.6 [+51.7,+169.4]	+104.5 [+41.3,+156.7]	-73.2	10/10	80.2	29.3	74.2
2:1	Conditional	+144.3	+27.7 [-23.9,+79.3]	+38.3 [+0.0,+54.4]	-116.7	7/10	74.6	66.7	41.6
2:1	All-row	+153.4	+30.8 [-12.3,+74.0]	+25.2 [+4.8,+52.2]	-122.5	7/10	63.7	55.4	54.1
2:1	Two-stage	+136.8	+37.8 [-2.7,+78.3]	+23.5 [+1.3,+61.8]	-99.0	7/10	44.8	61.9	48.6
5:1	Conditional	+96.2	-52.9 [-132.9,+27.1]	-17.4 [-31.3,+7.0]	-149.1	3/10	69.4	74.2	33.0
5:1	All-row	+104.7	-20.0 [-76.0,+36.0]	-4.4 [-17.0,+0.0]	-124.7	2/10	24.4	84.7	26.4
5:1	Two-stage	+81.9	-33.6 [-120.4,+53.2]	-2.6 [-13.9,+25.2]	-115.5	5/10	13.1	84.5	26.2
10:1	Conditional	+79.1	-81.7 [-194.0,+30.7]	-25.2 [-51.4,-7.4]	-160.7	0/10	54.3	91.7	12.3
10:1	All-row	+88.1	-95.9 [-222.7,+30.9]	-48.7 [-93.1,+0.0]	-184.0	1/10	10.6	89.4	17.4
10:1	Two-stage	+72.8	-86.0 [-230.7,+58.7]	-12.2 [-74.9,+1.3]	-158.8	3/10	7.8	90.5	15.3

4.4 Severity, affected-group, mechanism and batch transport

Selector transport was evaluated without changing endpoint architecture, endpoint training data or the fixed test realization. A scale-3 conditional selector applied at test scales 1, 2 and 3 yielded equal-cost utilities +4.0, +12.5 and +15.8, respectively. Conversely, a scale-1 selector applied at scale 3 yielded +5.9. Across the full 3×3 calibration-to-test matrix, conditional-router utility ranged from +2.1 to +17.9 per 10,000 (Fig. 4a). Positive means therefore did not imply severity-invariant thresholds or event rates.

Calibration on affected group 1 also transferred unevenly. Conditional Lite-CF prevention/retention changed from 94.7%/9.7% on group 1 to 80.8%/17.8%, 93.0%/21.1% and 90.2%/17.2% on groups 2-4. Equal-cost utility was +15.8, +6.1, +60.3 and +31.0, but the group-2 fitted-pair interval was -16.3 to +28.4. The threshold is thus not transferable merely because sensor groups share a representation size.

Leave-one-mechanism-out conditional calibration produced mean utilities +27.9, +10.4, +1.4 and +15.3 on Gaussian, offset, drift and stuck-at tests, respectively. Every fitted-pair interval included zero; the drift range was -41.8 to +20.9. Gaussian-only calibration and 70%-dominant calibration mixtures likewise changed mechanism-specific prevention, retention and utility. These audits motivate the proposed future support gate rather than a universal selector.

Batch composition materially changed the aggregate. The 91/165/1,180 strict observations in Batches 8/9/10 make the observation-weighted result predominantly a Batch 10 result. Frozen Lite-CF macro-AUROC was 0.8151 overall, 0.8574 under equal batch weight, 0.8022 in the worst batch, 0.8747 in pooled Batches 8-9 and 0.8022 in Batch 10 (Table 10). Support-gated equal-cost utility was +7.5 overall but -2.5 in the worst batch. The complete mechanism $\times$ batch $\times$ method file reports opportunities and all seven requested outcomes for each fitted pair.




Figure 3: **Selective recovery as an explicit decision policy.** **a**, descriptive cell-stratified test ranking curves averaged over 40 mechanism-fitted-pair cells; these are not deployable operating policies, whereas marked comparison points use calibration-frozen thresholds. **b**, always-base, always-recovery, frozen conditional, retrospectively support-gated and all-row safety policies. Utilities use equal correction and harm values. **c**, utility under fixed-safety and calibration-selected thresholds. The symmetric-log axis is linear from -25 to $+25$ utility units; exact values are retained in Table 9 and the source-data table. **d**, row-normalized out-of-fold confusion matrix for the all-row multinomial outcome. Counts pool the 10 calibration fits.




Figure 4: **Selector transport and acquisition-batch robustness.** **a**, conditional-router equal-cost utility for calibration and test severities 1–3. **b**, prevention (solid) and retention (dashed) after calibrating on sensor group 1 and testing groups 1–4. **c**, macro-AUROC under observation-weighted, equal-batch, worst-batch and Batches 8–9/Batch 10 views. **d**, leave-one-mechanism-out conditional-router utility; bars are means and whiskers are 95% fitted-pair intervals across 10 endpoint fits on fixed test rows.

Table 10: **Batch-aware strict-set macro-AUROC.** Overall values pool observations within each fixed mechanism–fitted-pair cell; equal-batch values give Batches 8, 9 and 10 equal weight; worst batch is the minimum batch mean. Batches 8–9 are pooled and compared with Batch 10. These are summaries of one fixed instrument, not temporal-population intervals.

Policy	Overall	Equal-batch	Worst batch	Batches 8–9	Batch 10
Always PDRF	0.8115	0.8528	0.7990	0.8689	0.7990
Always RO-PDRF-Lite	0.8175	0.8607	0.8044	0.8780	0.8044
Frozen Lite-CF	0.8151	0.8574	0.8022	0.8747	0.8022
Retrospective support gate	0.8132	0.8550	0.8005	0.8721	0.8005
All-row utility 1:1	0.8160	0.8591	0.8029	0.8771	0.8029

4.5 Sparse calibration limited selector stability

The formal Lite-CF calibration sets contained a mean of 1,861 rows, but only 89.7 were correctness disagreements: 36.1 favoured PDRF and 53.6 favoured RO-PDRF-Lite. Frozen thresholds averaged 0.710 and ranged from 0.55 to 0.90. Across 10 repeated grouped fits per endpoint pair, the arithmetic mean of the 100 individual seed-repeat AUROCs—the *mean-of-repeats AUROC*—was 0.720 (standard deviation 0.138; range 0.327–0.954), and thresholds ranged from 0.525 to 0.90. This aggregation differs from the *pooled-row AUROC* of 0.562 computed by concatenating the formal out-of-fold predictions across informative calibration rows. Only 5/10 pairs passed the post hoc gate under the conservative reference criteria. The 94.7%/9.7% result remains the auditable ungated estimate; the retrospectively gated 96.2%/6.6% result illustrates the intended prospective fallback rule.

Coefficient signs were most stable for Lite confidence (94%), base confidence (89%) and confidence change (83%). Dominant-sign fractions for the entropy terms were 61–67%. Strong shrinkage traded away conditional discrimination: mean-of-repeats AUROC increased from 0.690 at $C = 0.01$ to 0.717 at $C = 0.05$, 0.800 at $C = 1$ and 0.811 at $C = 5$. The original router setting $C = 0.05$ was selected before the post hoc gate was proposed; it is a conservative estimation choice, not the value that maximizes calibration preference AUROC.

Calibration-size sensitivity showed the same boundary. Using 25%, 50%, 75% and 100% of calibration groups produced pooled mean prevention of 88.7%, 91.6%, 92.4% and 93.9%, respectively. Retention changed from 12.1% to 10.0%, while macro-AUROC remained near 0.81. These are repeated-subsampling summaries and differ from the equal-weight 40-cell headline estimand.

Endpoint temperature scaling changed the conditional router: uncalibrated endpoint features gave pooled-row preference AUROC 0.563, prevention 95.7%, retention 6.6% and utility +10.8, compared with pooled-row AUROC 0.562, 94.7%, 9.7% and +15.8 after temperature scaling. Platt calibration of the conditional score reduced preference NLL from 0.694 to 0.689 and Brier score from 0.249 to 0.239 without changing decisions when the frozen threshold was monotonically mapped. Isotonic calibration introduced flat regions and changed decisions, increasing utility but reducing prevention to 82.1%; it is therefore a sensitivity, not a replacement. Probability calibration and hard-decision selection remain separate objectives [34, 39].




Figure 5: **Selector stability and final probability quality.** **a**, seed-wise out-of-fold conditional-preference AUROC and threshold distributions across 100 repeated grouped fits. **b**, median standardized Lite-CF coefficient with interquartile range; percentages are dominant-sign frequencies. **c**, reliability curves lie predominantly below the identity line, showing substantial overconfidence for all three outputs. **d**, negative log-likelihood and 15-bin expected calibration error after separating endpoint-class agreement from disagreement. Lite-CF emits one endpoint vector, not a mixture; score-calibration controls are reported separately.

4.6 Hydraulic single-rig sensitivity boundary

The guarded hydraulic analysis used four native component-condition tasks and controlled pressure- and vibration-group faults (Table 11). RO-PDRF-Full exceeded recovery-objective-matched RO-CAGF in six of eight task-fault conditions. The mean difference was +0.015, but the task-fault-seed hierarchical interval (-0.006 to $+0.036$) included zero. Performance ranged from near ceiling for cooler (RO-PDRF-Full macro-AUROC 0.996) to near chance for pump and accumulator (0.528 and 0.487), limiting architecture interpretation on the latter tasks.

The earlier random-stratified split yielded a larger mean difference (+0.044), indicating optimistic adjacent-cycle allocation. Guard bands reduced direct adjacency but did not create independent devices. No-interpolation and random-split results are therefore supplementary sensitivity evidence; the main result supports single-rig platform-internal robustness, not a general architecture ranking.

Table 11: **Hydraulic single-rig sensitivity analysis.** Guarded-block affected-subset macro-AUROC values are means $\pm$ standard deviations across five matched seeds. The eight task-fault rows are conditions from one rig, not independent devices; the earlier random-stratified results are supplementary sensitivity evidence.

Task	Fault	RO-CAGF	RO-PDRF-Full	Difference
Cooler	Pressure	0.942 ± 0.023	0.992 ± 0.004	+0.050
Cooler	Vibration	0.947 ± 0.029	1.000 ± 0.000	+0.053
Valve	Pressure	0.730 ± 0.026	0.762 ± 0.023	+0.032
Valve	Vibration	0.739 ± 0.029	0.749 ± 0.032	+0.010
Pump	Pressure	0.569 ± 0.035	0.536 ± 0.018	-0.033
Pump	Vibration	0.533 ± 0.061	0.521 ± 0.022	-0.012
Accumulator	Pressure	0.478 ± 0.022	0.491 ± 0.023	+0.013
Accumulator	Vibration	0.474 ± 0.047	0.483 ± 0.044	+0.009

4.7 AHU direct-diagnosis boundary

The chronological air-handling-unit tests contained 37,831 untouched future-time records, including 2,069 expert-annotated temperature-sensor faults. Because the target is the fault itself and no independent downstream state is available, this experiment cannot test whether suppressing that sensor recovers another prediction. Early fusion exceeded bounded recovery in all 15 building-seed pairs (mean AUROC difference -0.0341 for bounded recovery, $P < 0.001$), and leave-one-building-out bounded recovery collapsed to 0.298 ± 0.005 in the hospital. The operational data therefore establish a direct-diagnosis and transfer boundary, not downstream recovery validation; complete subtype, shift, adaptation and risk analyses remain in the Supplementary Information.

4.8 Routed probability quality depended on endpoint agreement

Final routed probabilities were evaluated directly rather than inferred from the residual-harm rate. Reliability curves lay predominantly below the identity line, indicating substantial overconfidence for all methods. On strict endpoint-class agreements, Lite-CF NLL/ECE were 2.667/0.342, between PDRF (2.821/0.349) and RO-PDRF-Lite (2.601/0.342). On disagreements, Lite-CF NLL/ECE

were 2.813/0.344; unconditional Lite was better at 2.616/0.207, whereas PDRF was worse at 2.869/0.376. The software can expose the selected endpoint identity and conditional preference score with each probability vector. Downstream controllers should therefore treat confidence as endpoint-specific, not as one globally calibrated routed probability (Fig. 5c,d). Post-routing recalibration remains future work.

4.9 Lite-CF reduced, but did not remove, deployment cost

PDRF and RO-PDRF-Lite each require one forward pass, 19,740 parameters and 38,528 counted linear-layer FLOPs per observation. Lite-CF requires two passes, 39,487 parameters and 77,056 endpoint FLOPs. With 12 intra-op threads, mean seed-level median batch-one latencies were 0.602 ± 0.263 ms for PDRF, 0.589 ± 0.258 ms for RO-PDRF-Lite and 2.471 ± 0.865 ms for Lite-CF; complete-batch latencies were 4.12 ± 0.73 , 3.65 ± 0.66 and 10.21 ± 2.09 ms. At one intra-op thread, the corresponding batch-one values were 0.273 ± 0.090 , 0.252 ± 0.068 and 1.553 ± 0.447 ms, and complete-batch values were 9.35 ± 1.77 , 5.17 ± 0.58 and 11.60 ± 0.81 ms. All measurements used the same Intel Core i5-13500H, PyTorch 2.12.1+cpu, three warm-ups and seven timed calls, excluding loading. They do not establish embedded latency, energy efficiency or real-time readiness.

Event-normalized engineering cost was large. Across the 40 strict cells, the frozen router required 57,440 extra endpoint passes to produce 139 retained corrections and 91 net corrections: 413 extra passes per retained correction and 631 per net correction. It selected recovery 29,970 times, corresponding to 216 selected rows per retained correction and 329 per net correction. The support gate reduced total extra passes but also reduced net corrections, yielding 668 extra passes and 416 selected rows per net correction. For batch-one timing, the routed call exceeded the sum of separately timed endpoints by 1.03 ms at one thread and 1.28 ms at 12 threads; this residual includes feature construction, logistic scoring, array selection and Python orchestration. Full-batch residuals were non-additive (-2.93 and $+2.45$ ms), so they are reported as orchestration/timing residuals rather than pure router latency.

5 Discussion

The strict estimand and frozen-run discipline narrow the central interpretation. Clean-to-faulted distillation produced a useful recovery endpoint, but the bounded response architecture alone did not solve deployment choice. The stronger contribution is a decision protocol that distinguishes fault assignment from an intervention that reaches an available group, measures correction and error transfer, separates safety from utility and freezes every operating threshold before test evaluation. It should not be read as a general theory of sensor reliability.

Cross-fitting addresses threshold leakage, not selected-sample bias or sparse-event estimation. The original conditional score used only 89.7 correctness disagreements on average and therefore estimated preference conditional on disagreement. The all-row multinomial and two-stage controls target the required deployment population, but their rare base-better and recovery-better classes remained difficult. No test curve dominated the frozen point in both prevention and retention, while class-conditional and margin rules obtained higher equal-cost utility only by accepting substantially more harm. The estimand correction therefore changes what is measurable without creating a uniformly superior router.

The post hoc proposed support gate is central to the intended prospective deployment protocol, but the present evidence is retrospective. Under the conservative reference configuration only 5/10 endpoint pairs passed, and retrospective fallback of the other half to PDRF achieved 96.2% prevention and 6.6% retention. Its fitted-pair utility interval crossed zero and worst-batch utility

was -2.5 per 10,000. Across 27 gate configurations, passed pairs ranged from one to seven and decision outcomes changed accordingly. A future site should freeze its support rule before outcome evaluation, avoid transporting coefficients or thresholds from another fitted pair, and use the base endpoint until local support is adequate.

Utility-aware thresholds clarify when selection adds value. At equal event values, unconditional recovery had the largest mean utility. On this fixed test set and with no compute penalty, two-stage routing had the highest mean at a two-to-one harm ratio, but its fitted-pair interval crossed zero and calibration-to-test utility decreased by 99.0 units per 10,000. At ratios of five and ten, always PDRF dominated every learned policy by mean utility. No single “safe” threshold is therefore application-independent. The values V_C , C_H and λ must be defined in operational units—for example, avoided inspection cost, cost of a wrong control decision and compute/latency cost converted to the same per-10,000-observation scale—before calibration. A monotone score calibration cannot substitute for that value specification. The all-row and two-stage variants remain analytical controls because this study did not define independent support gates for them; failure of the conditional gate cannot authorize switching to either variant.

The two-pass design avoids leave-one-group-out inference but stores two 19,740-parameter endpoints (154.25 KiB in float32) and evaluates both. Pooled over the strict audit, frozen Lite-CF required approximately 413 extra passes per retained correction and 631 per net correction; about 329 recovery selections were required per net correction. The support gate required 668 extra passes and 416 selected rows per net correction because unsupported fits contributed no routed gains. Batch-one orchestration residual beyond the two separately timed endpoints was 1.03 ms at one thread and 1.28 ms at 12 threads; the full-batch residual ranged from -2.93 to $+2.45$ ms because independent timings are non-additive. These workstation measurements are not embedded latency or energy guarantees.

Mechanism, severity, affected-group and acquisition-batch transport all changed the decision balance. Leave-one-mechanism-out fitting intervals included zero; group-2 prevention fell to 80.8%; Batch 10 supplied 82.2% of strict observations and pulled the observation-weighted AUROC toward its lower value. These analyses are more informative than one aggregate mean, but they still reuse one chemical instrument and fixed rows. The fitted-pair intervals quantify optimization variation, not sampling from devices or natural faults.

Transfer tests define boundaries rather than confirm field recovery. The hydraulic result is single-rig sensitivity, and its guarded architecture interval crosses zero. The AHU target is the sensor fault itself, so diagnosis may reward signatures that recovery suppresses. A prospective study must include multiple independently operated devices, maintenance-adjudicated sensor failures and downstream targets that do not encode those failures. Endpoint fitting and site-specific calibration must precede outcome analysis, with the support gate, PDRF fallback and utility values frozen in advance.

Further boundaries remain. Routed probabilities are overconfident and endpoint-specific; no tested output dominates safety, ranking and calibration. The present utility treats corrections and harms as class-independent binary events. It does not represent asymmetric class-to-class misclassification costs or differences between two incorrect endpoint predictions. Response saturation prevents device-health interpretation, and erroneous clean teachers or quality metadata can transfer errors. Evidence remains limited to controlled feature corruption, one chemical instrument, one hydraulic rig and expert/rule-based AHU labels. Multi-device prospective evaluation is required before field-recovery claims.

6 Conclusion

Support-gated selective recovery converted a useful but fallible recovery endpoint into an explicit, auditable risk-control protocol. The frozen ungated router prevented 94.7% of negative transfers while retaining 9.7% of corrections. A post hoc proposed gate, evaluated retrospectively on the frozen experiment, admitted only 5/10 fitted pairs under its conservative reference configuration; fallback of the remainder to PDRF yielded 96.2% prevention and 6.6% retention. All-row and two-stage analytical controls corrected the conditional-estimand mismatch, and two-stage had the highest mean utility among tested policies at a two-to-one harm ratio on this fixed test set, but its fitted-pair interval crossed zero. Always recovery was preferred at equal event values, whereas always PDRF dominated by mean utility at ratios of five or ten. Severity, sensor-group, unseen-mechanism and batch audits showed that thresholds did not transport uniformly. The contribution is a reproducible evaluation and future deployment protocol for controlled available-group corruption, not prospective validation, calibrated device health, natural-failure recovery or device-independent field performance.

Data and code availability

The Gas Sensor Array Drift Dataset (DOI: <https://doi.org/10.24432/C5RP6W>) and Condition Monitoring of Hydraulic Systems dataset (DOI: <https://doi.org/10.24432/C5CW21>) are available from the UCI Machine Learning Repository under CC BY 4.0. The operational AHU data are available from Figshare under CC BY 4.0 (DOI: <https://doi.org/10.6084/m9.figshare.27147678.v3>). Code, fixed seeds, environment information, prediction-level source data and figure/table scripts are publicly available in the `cee_revision` directory at <https://github.com/Coreey9999/recovery-supervision-calibrated-fallback>. The exact computational snapshot for the support-gate, fitted-pair utility, all-row, transport and batch audits is commit `59be82fc18aaf48b819c3c71a9482425f8baac45`; the versioned submission release is `v1.5.0-round2-support-audit`. Exact dependencies are recorded in `requirements-lock.txt` and `environment.yml`, and the script-to-result manifest is Supplementary Table S1. A permanent software DOI has not been assigned.

Ethics statement

The study used public engineering datasets and controlled computer-generated sensor corruptions. It involved no human participants, personal data, animals or biological specimens; institutional ethics review was not applicable.

Author contributions

Riyang Luo: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Visualization, Writing—original draft, Writing—review and editing, and Project administration.

Acknowledgements and funding

The author received no specific funding for this work and has no acknowledgements to declare.

Declaration of interests

The author declares that there are no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

References

- [1] X. Liu, J. Zhang, S. Zhou, et al. Towards deployment-centric multimodal AI beyond vision and language. *Nature Machine Intelligence*, 7:1612–1624, 2025. doi: 10.1038/s42256-025-01116-5.
- [2] Tapabrata Chakraborti, Christopher R. S. Banerji, Ariane Marandon, et al. Personalized uncertainty quantification in artificial intelligence. *Nature Machine Intelligence*, 7:522–530, 2025. doi: 10.1038/s42256-025-01024-8.
- [3] Robin Mitra, Sarah F. McGough, Tapabrata Chakraborti, et al. Learning from data with structured missingness. *Nature Machine Intelligence*, 5:13–23, 2023. doi: 10.1038/s42256-022-00596-z.
- [4] Renjie Wu, Hu Wang, Hsiang-Ting Chen, and Gustavo Carneiro. Deep multimodal learning with missing modality: A survey. *Transactions on Machine Learning Research*, February 2026. URL <https://openreview.net/forum?id=tc7RFcx4hT>.
- [5] Natalia Neverova, Christian Wolf, Graham Taylor, and Florian Nebout. ModDrop: Adaptive multi-modal gesture recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(8):1692–1706, 2016. doi: 10.1109/TPAMI.2015.2461544.
- [6] Mingcheng Li, Dingkan Yang, Yuxuan Lei, Shunli Wang, Shuaibing Wang, Liuzhen Su, Kun Yang, Yuzheng Wang, Mingyang Sun, and Lihua Zhang. A unified self-distillation framework for multimodal sentiment analysis with uncertain missing modalities. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 10074–10082, 2024. doi: 10.1609/aaai.v38i9.28871.
- [7] Sangmin Woo, Sumin Lee, Yeonju Park, Muhammad Adi Nugroho, and Changick Kim. Towards good practices for missing modality robust action recognition. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 2776–2784, 2023. doi: 10.1609/aaai.v37i3.25378.
- [8] Hao Wang, Shengda Luo, Guosheng Hu, and Jianguo Zhang. Gradient-guided modality decoupling for missing-modality robustness. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 15483–15491, 2024. doi: 10.1609/aaai.v38i14.29474.
- [9] Brandon Mckinzie, Vaishaal Shankar, Joseph Yitan Cheng, Yinfei Yang, Jonathon Shlens, and Alexander T. Toshev. Robustness in multimodal learning under train–test modality mismatch. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 24291–24303. PMLR, 2023. URL <https://proceedings.mlr.press/v202/mckinzie23a.html>.
- [10] Wenfang Yao, Kejing Yin, William K. Cheung, Jia Liu, and Jing Qin. DrFuse: Learning disentangled representation for clinical multi-modal fusion with missing modality and modal inconsistency. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 16416–16424, 2024. doi: 10.1609/aaai.v38i15.29578.

- [11] Wuliang Huang, Yiqiang Chen, Xinlong Jiang, Chenlong Gao, Teng Zhang, Qian Chen, and Yifan Wang. Mitigating pervasive modality absence through multimodal generalization and refinement. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 26796–26804, 2025. doi: 10.1609/aaai.v39i25.34883.
- [12] Zixian Gao, Xun Jiang, Xing Xu, Fumin Shen, Yujie Li, and Heng Tao Shen. Embracing unimodal aleatoric uncertainty for robust multimodal fusion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 26876–26885, 2024. doi: 10.1109/CVPR52733.2024.02538.
- [13] Minkyung Cho, Yulong Cao, Jiachen Sun, Qingzhao Zhang, Marco Pavone, Jeong Joon Park, Heng Yang, and Zhuoqing Mao. Cocoon: Robust multi-modal perception with uncertainty-aware sensor fusion. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/6580845ce91a748a9bb09cec079aac1d-Abstract-Conference.html.
- [14] Konyul Park, Yecheol Kim, Daehun Kim, and Jun Won Choi. Resilient sensor fusion under adverse sensor failures via multi-modal expert fusion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6720–6729, 2025. doi: 10.1109/CVPR52734.2025.00630.
- [15] Sahil Sidheekh, Pranuthi Tenali, Saurabh Mathur, Erik Blasch, Kristian Kersting, and Sriraam Natarajan. Credibility-aware multimodal fusion using probabilistic circuits. In *Proceedings of the 28th International Conference on Artificial Intelligence and Statistics*, volume 258 of *Proceedings of Machine Learning Research*, pages 2305–2313. PMLR, 2025. URL <https://proceedings.mlr.press/v258/sidheekh25a.html>.
- [16] Wei Liu, Yufei Chen, and Xiaodong Yue. Enhancing multi-view classification reliability with adaptive rejection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 18969–18977, 2025. doi: 10.1609/aaai.v39i18.34088.
- [17] Stephan Rabanser, Anvith Thudi, Kimia Hamidieh, Adam Dziedzic, Israfil Bahceci, Akram Bin Sediq, Hamza Sokun, and Nicolas Papernot. Selective prediction via training dynamics. *Transactions on Machine Learning Research*, 2025. URL <https://openreview.net/forum?id=2wgnepQjyF>.
- [18] Qingyang Zhang, Haitao Wu, Changqing Zhang, Qinghua Hu, Huazhu Fu, Joey Tianyi Zhou, and Xi Peng. Provable dynamic fusion for low-quality multimodal data. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 41753–41769. PMLR, 2023. URL <https://proceedings.mlr.press/v202/zhang23ar.html>.
- [19] Jyrki Kullaa. Distinguishing between sensor fault, structural damage, and environmental or operational effects in structural health monitoring. *Mechanical Systems and Signal Processing*, 25(8):2976–2989, 2011. doi: 10.1016/j.ymssp.2011.05.017.
- [20] Jyrki Kullaa. Detection, identification, and quantification of sensor fault in a sensor network. *Mechanical Systems and Signal Processing*, 40(1):208–221, 2013. doi: 10.1016/j.ymssp.2013.05.007.

- [21] Nikolai Helwig, Eliseo Pignanelli, and Andreas Schütze. Detecting and compensating sensor faults in a hydraulic condition monitoring system. In *Proceedings SENSOR 2015*, pages 641–646, 2015. doi: 10.5162/sensor2015/D8.1.
- [22] Nikolai Helwig, Eliseo Pignanelli, and Andreas Schütze. Condition monitoring of a complex hydraulic system using multivariate statistics. In *2015 IEEE International Instrumentation and Measurement Technology Conference Proceedings*, pages 210–215, 2015. doi: 10.1109/I2MTC.2015.7151267.
- [23] Xiangkai Ma, Pei Wang, Bozhou Zhang, and Ming Sun. A multirate sensor information fusion strategy for multitask fault diagnosis based on convolutional neural network. *Journal of Sensors*, 2021:9952450, 2021. doi: 10.1155/2021/9952450.
- [24] Aurelien Teguede Keleko, Bernard Kamsu-Foguem, Raymond Houe Ngouna, and Amèvi Tongne. Health condition monitoring of a complex hydraulic system using deep neural network and DeepSHAP explainable XAI. *Advances in Engineering Software*, 175:103339, 2023. doi: 10.1016/j.advensoft.2022.103339.
- [25] Cevahir Yildirim, Alba M. Franco-Pereira, and Rosa E. Lillo. Condition monitoring and multi-fault classification of hydraulic systems using multivariate functional data analysis. *Heliyon*, 11(1):e41251, 2025. doi: 10.1016/j.heliyon.2024.e41251.
- [26] Seunghyeon Wang. Real operational labeled data of air handling units from office, auditorium, and hospital buildings. *Scientific Data*, 12:1481, 2025. doi: 10.1038/s41597-025-05825-9.
- [27] Seunghyeon Wang. Labelled dataset for FDD in different building types: auditorium, hospital, and large-scale office. Dataset, version 3, 2024.
- [28] Seunghyeon Wang. Automated fault diagnosis detection of air handling units using real operational labelled data and transformer-based methods at 24-hour operation hospital. *Building and Environment*, 282:113257, 2025. doi: 10.1016/j.buildenv.2025.113257.
- [29] Bowei Feng, Qizhen Zhou, Jianchun Xing, Qiliang Yang, Yin Chen, and Zhongkai Deng. Attention-empowered transfer learning method for HVAC sensor fault diagnosis in dynamic building environments. *Building and Environment*, 250:111148, 2024. doi: 10.1016/j.buildenv.2023.111148.
- [30] Seunghyeon Wang. Temporal-contextual self-supervised time-series learning for automated fault detection and diagnosis of air handling units in buildings. *Building and Environment*, 292:114300, 2026. doi: 10.1016/j.buildenv.2026.114300.
- [31] Alex Kendall and Yarin Gal. What uncertainties do we need in bayesian deep learning for computer vision? In *Advances in Neural Information Processing Systems*, volume 30, 2017. URL <https://proceedings.neurips.cc/paper/2017/hash/2650d6089a6d640c5e85b2b88265dc2b-Abstract.html>.
- [32] Eyke Hüllermeier and Willem Waegeman. Aleatoric and epistemic uncertainty in machine learning: An introduction to concepts and methods. *Machine Learning*, 110:457–506, 2021. doi: 10.1007/s10994-021-05946-3.

- [33] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems*, volume 30, 2017. URL <https://proceedings.neurips.cc/paper/2017/hash/9ef2ed4b7fd2c810847ffa5fa85bce38-Abstract.html>.
- [34] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR, 2017. URL <https://proceedings.mlr.press/v70/guo17a.html>.
- [35] Yaniv Ovadia, Emily Fertig, Jie Ren, Zachary Nado, D. Sculley, Sebastian Nowozin, Joshua V. Dillon, Balaji Lakshminarayanan, and Jasper Snoek. Can you trust your model’s uncertainty? evaluating predictive uncertainty under dataset shift. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://proceedings.neurips.cc/paper/2019/hash/8558cb408c1d76621371888657d2eb1d-Abstract.html>.
- [36] Matthias Minderer, Josip Djolonga, Rob Romijnders, Frances Hubis, Xiaohua Zhai, Neil Houlsby, Dustin Tran, and Mario Lucic. Revisiting the calibration of modern neural networks. In *Advances in Neural Information Processing Systems*, volume 34, pages 15682–15694, 2021. URL <https://proceedings.neurips.cc/paper/2021/hash/8420d359404024567b5aefda1231af24-Abstract.html>.
- [37] Juozas Vaicenavicius, David Widmann, Carl Andersson, Fredrik Lindsten, Jacob Roll, and Thomas B. Schön. Evaluating model calibration in classification. In *Proceedings of the Twenty-Second International Conference on Artificial Intelligence and Statistics*, volume 89 of *Proceedings of Machine Learning Research*, pages 3459–3467, 2019. URL <https://proceedings.mlr.press/v89/vaicenavicius19a.html>.
- [38] Roderick J. A. Little and Donald B. Rubin. *Statistical Analysis with Missing Data*. Wiley, 3 edition, 2019. doi: 10.1002/9781119482260.
- [39] Meelis Kull, Miquel Perello-Nieto, Markus Kängsepp, Telmo Silva Filho, Hao Song, and Peter Flach. Beyond temperature scaling: Obtaining well-calibrated multi-class probabilities with dirichlet calibration. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://proceedings.neurips.cc/paper/2019/hash/8ca01ea920679a0fe3728441494041b9-Abstract.html>.
- [40] Thomas Mortier, Viktor Bengs, Eyke Hullermeier, Stijn Luca, and Willem Waegeman. On the calibration of probabilistic classifier sets. In *Proceedings of the 26th International Conference on Artificial Intelligence and Statistics*, volume 206 of *Proceedings of Machine Learning Research*, pages 8857–8870, 2023. URL <https://proceedings.mlr.press/v206/mortier23a.html>.
- [41] Aleksandr Podkopaev and Aaditya Ramdas. Distribution-free uncertainty quantification for classification under label shift. In *Proceedings of the Thirty-Seventh Conference on Uncertainty in Artificial Intelligence*, volume 161 of *Proceedings of Machine Learning Research*, pages 844–853, 2021. URL <https://proceedings.mlr.press/v161/podkopaev21a.html>.
- [42] Baochen Sun, Jiashi Feng, and Kate Saenko. Return of frustratingly easy domain adaptation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 30, pages 2058–2065, 2016. doi: 10.1609/aaai.v30i1.10306.

- [43] Shiori Sagawa, Pang Wei Koh, Tatsunori B. Hashimoto, and Percy Liang. Distributionally robust neural networks for group shifts: On the importance of regularization for worst-case generalization. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=ryxGuJrFvS>.
- [44] M. Jehanzeb Mirza, Jakub Micorek, Horst Possegger, and Horst Bischof. The norm must go on: Dynamic unsupervised domain adaptation by normalization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14765–14775, 2022. URL https://openaccess.thecvf.com/content/CVPR2022/html/Mirza_The_Norm_Must_Go_On_Dynamic_Unsupervised_Domain_Adaptation_by_CVPR_2022_paper.html.
- [45] Ran El-Yaniv and Yair Wiener. On the foundations of noise-free selective classification. *Journal of Machine Learning Research*, 11(53):1605–1641, 2010. URL <https://jmlr.org/papers/v11/el-yaniv10a.html>.
- [46] Yonatan Geifman and Ran El-Yaniv. SelectiveNet: A deep neural network with an integrated reject option. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2151–2159. PMLR, 2019. URL <https://proceedings.mlr.press/v97/geifman19a.html>.
- [47] Vojtech Franc, Daniel Prusa, and Vaclav Voracek. Optimal strategies for reject option classifiers. *Journal of Machine Learning Research*, 24(11):1–49, 2023. URL <https://www.jmlr.org/papers/v24/21-0048.html>.
- [48] Hengyue Liang, Le Peng, and Ju Sun. Selective classification under distribution shifts. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=dmxMGW6J7N>.
- [49] Hussein Mozannar and David Sontag. Consistent estimators for learning to defer to an expert. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 7076–7087. PMLR, 2020. URL <https://proceedings.mlr.press/v119/mozannar20b.html>.
- [50] Dharmesh Tailor, Aditya Patra, Rajeev Verma, Putra Manggala, and Eric Nalisnick. Learning to defer to a population: A meta-learning approach. In *Proceedings of the 27th International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings of Machine Learning Research*, pages 3475–3483. PMLR, 2024. URL <https://proceedings.mlr.press/v238/tailor24a.html>.
- [51] Minmin Chen, Zhixiang Xu, Kilian Weinberger, Olivier Chapelle, and Dor Kedem. Classifier cascade for minimizing feature evaluation cost. In *Proceedings of the Fifteenth International Conference on Artificial Intelligence and Statistics*, volume 22 of *Proceedings of Machine Learning Research*, pages 218–226. PMLR, 2012. URL <https://proceedings.mlr.press/v22/chen12c.html>.
- [52] Rafael M. O. Cruz, Robert Sabourin, and George D. C. Cavalcanti. Dynamic classifier selection: Recent advances and perspectives. *Information Fusion*, 41:195–216, 2018. doi: 10.1016/j.inffus.2017.09.010.
- [53] Thomas G. Dietterich. Approximate statistical tests for comparing supervised classification learning algorithms. *Neural Computation*, 10(7):1895–1923, 1998. doi: 10.1162/089976698300017197.

- [54] Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30, 2006. URL <https://www.jmlr.org/papers/v7/demsar06a.html>.
- [55] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1994. doi: 10.1201/9780429246593.
- [56] Quinn McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947. doi: 10.1007/BF02295996.
- [57] Takaya Saito and Marc Rehmsmeier. The precision–recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets. *PLOS ONE*, 10(3):e0118432, 2015. doi: 10.1371/journal.pone.0118432.
- [58] Antti Tarvainen and Harri Valpola. Mean teachers are better role models: Weight-averaged consistency targets improve semi-supervised deep learning results. In *Advances in Neural Information Processing Systems*, volume 30, pages 1195–1204, 2017. URL https://proceedings.neurips.cc/paper_files/paper/2017/hash/68053af2923e00204c3ca7c6a3150cf7-Abstract.html.
- [59] Zhenhuan Yang, Yunwen Lei, Siwei Lyu, and Yiming Ying. Stability and differential privacy of stochastic gradient descent for pairwise learning with non-smooth loss. In *Proceedings of the 24th International Conference on Artificial Intelligence and Statistics*, volume 130 of *Proceedings of Machine Learning Research*, pages 2026–2034, 2021. URL <https://proceedings.mlr.press/v130/yang21c.html>.
- [60] Alexander Vergara. Gas sensor array drift dataset, 2012. Licensed under CC BY 4.0.
- [61] Alexander Vergara, Shankar Vembu, Tuba Ayhan, Margaret A. Ryan, Margie L. Homer, and Ramon Huerta. Chemical gas sensor drift compensation using classifier ensembles. *Sensors and Actuators B: Chemical*, 166–167:320–329, 2012. doi: 10.1016/j.snb.2012.01.074.
- [62] Nikolai Helwig, Eliseo Pignanelli, and Andreas Schütze. Condition monitoring of hydraulic systems, 2018. Licensed under CC BY 4.0.